

Biostatistics 615 - Statistical Computing

Lecture 12 Advanced R – Performance

Jian Kang

Nov 3, 2015

Microbenchmarking

Accurate Timing Functions

- A microbenchmark is a measurement of the performance of a very small piece of code, something that might take microseconds (μs) or nanoseconds (ns) to run
- It provides very precise timings, making it possible to compare operations that only take a tiny amount of time. For example, the following code compares the speed of three ways of computing a square root.

```
> library(microbenchmark)
> x = runif(1000)
> microbenchmark(sqrt(x), x^0.5, exp(0.5*log(x)))
```

Unit: microseconds

expr	min	lq	mean	median	uq	max	neval	cld
sqrt(x)	4.379	4.5280	7.62580	7.3755	7.5765	71.789	100	a
x^0.5	26.990	27.2045	29.20157	30.1425	30.3490	42.836	100	c
exp(0.5 * log(x))	15.196	15.4210	17.61316	18.4675	18.7420	29.355	100	b

Example: for loop versus build-in function

```
cumsum_1 = function(x){  
  for(i in 2:length(x)){  
    x[i] = x[i] + x[i-1]  
  }  
  return(x)  
}
```

```
cumsum_2 = function(x){  
  sapply(2:length(x), function(i) x[i] <- x[i] + x[i-1] )  
  return(x)  
}
```

```
> x = rep(1,length=1000)
```

```
> microbenchmark(my_cumsum_1(x), my_cumsum_2(x), cumsum(x))
```

Unit: microseconds

	expr	min	lq	mean	median	uq	max	neval	cld
cumsum_1(x)	809.002	890.733	1024.36413	952.0435	1105.4545	2611.256	100	b	
cumsum_2(x)	1160.065	1289.626	1542.70321	1421.2905	1554.9715	3211.737	100	c	
cumsum(x)	2.269	3.182	3.83002	3.6445	4.1175	8.915	100	a	

Example: sapply versus apply

```
> x = matrix(rnorm(1e6),nrow=100,ncol=1e4)
>
> microbenchmark(
+   y1 <- sapply(1:1e4,function(i) cumsum(x[,i])),
+   y2 <- apply(x,2,cumsum)
+ )
```

Unit: milliseconds

			expr	min	lq	mean	median
y1 <- sapply(1:10000, function(i) cumsum(x[, i]))				37.81795	44.84013	71.25728	49.55951
y2 <- apply(x, 2, cumsum)				40.64251	53.01257	85.55411	61.81404
uq	max	neval	cld				
121.7896	150.2417	100	a				
132.0372	163.9973	100	b				

```
> all(y1==y2)
[1] TRUE
```

Better implementation ?

```
library(Rcpp)
```

```
cppFunction('void cumsum_3(NumericVector& x){  
    for(int i = 0; i<x.size(); i++){  
        x[i] = x[i] + x[i-1];  
    }  
}')  
)
```

by reference

```
> microbenchmark(cumsum_1(rep(1,length=1000)),  
+               cumsum_2(rep(1,length=1000)),  
+               cumsum_3(rep(1,length=1000)),  
+               cumsum(rep(1,length=1000)),times=1000)
```

Unit: microseconds

	expr	min	lq	mean	median
cumsum_1(rep(1, length = 1000))		766.112	875.300	1076.159666	929.0250
cumsum_2(rep(1, length = 1000))		1118.403	1250.622	1439.504496	1328.0435
cumsum_3(rep(1, length = 1000))		5.854	6.868	8.420184	7.7005
cumsum(rep(1, length = 1000))		3.677	4.739	5.399996	5.0330
uq	max	neval	cld		
998.760	76564.592	1000	b		
1424.347	3420.567	1000	c		
9.371	26.128	1000	a		
5.568	38.861	1000	a		

High Performance Functions With Rcpp

This magic comes by way of the Rcpp package, a fantastic tool makes it very simple to connect C++ to R

Typical bottlenecks that C++ can address include:

- Loops that can not be easily vectorised because subsequent iterations depend on previous ones.
- Recursive functions, or problems which involve calling functions millions of times. The overhead of calling a function in C++ is much lower than that in R
- Problems that require advanced data structures and algorithms that R does not provide. Through the standard template library (STL), C++ has efficient implementations of many important data structures, from ordered maps to double-ended queues.

Prerequisites

All examples in this lecture need version 0.10.1 or above of the `Rcpp` package. This version includes `cppFunction()` and `sourceCpp()`, which makes it very easy to connect `C++` to `R`. Install the latest version of `Rcpp` from CRAN with `install.packages("Rcpp")`. You'll also need a working `C++` compiler. To get it:

- On Windows, install Rtools
<https://cran.r-project.org/bin/windows/Rtools/>.
- On Mac, install Xcode from the app store.
- On Linux, `sudo apt-get install r-base-dev` or similar.

cppFunction()

It allows you to write C++ functions in R

```
library(Rcpp)

cppFunction('int add(int x, int y, int z) {
  int sum = x + y + z;
  return sum;
}')

# add works like a regular R function
add

#> function (x, y, z)
#> .Primitive(".Call")(<pointer: 0x7f2f4aa933d0>, x, y, z)
add(1, 2, 3)
#> [1] 6
```

When you run this code, Rcpp will compile the C++ code and construct an R function that connects to the compiled C++ function.

We will summarize the basics by translating simple R functions to their C++ equivalents. We start simple with a function that has no inputs and a scalar output, and then get progressively more complicated:

No inputs, scalar output

The R function is

```
one <- function() 1L
```

The equivalent C++ function is:

```
int one() {  
    return 1;  
}
```

We can compile and use this from R with `cppFunction()`

```
cppFunction('int one() {  
    return 1;  
}')
```

Difference between C++ and R functions

This small function illustrates a number of important differences between R and C++:

- The syntax to create a function looks like the syntax to call a function; you do not use assignment to create functions as you do in R.
- You must declare the type of output the function returns. This function returns an int (a scalar integer). The classes for the most common types of R vectors are: `NumericVector`, `IntegerVector`, `CharacterVector`, and `LogicalVector`.
- Scalars and vectors are different. The scalar equivalents of `numeric`, `integer`, `character`, and `logical` vectors are: `double`, `int`, `String`, and `bool`.
- You must use an explicit `return` statement to return a value from a function.
- Every statement is terminated by a `;`.

Scalar input, scalar output

A scalar version of the `sign()` function which returns `1` if the input is positive, and `-1` if it's negative:

```
signR <- function(x) {  
  if (x > 0) {  
    1  
  } else if (x == 0) {  
    0  
  } else {  
    -1  
  }  
}  
  
cppFunction('int signC(int x) {  
  if (x > 0) {  
    return 1;  
  } else if (x == 0) {  
    return 0;  
  } else {  
    return -1;  
  }  
}')
```

Vector input, scalar output

One big difference between **R** and **C++** is that the cost of loops is much lower in **C++**. For example, we could implement the sum function in **R** using a loop.

In **C++**, loops have very little overhead. In **STL**, you will see alternatives to for loops that more clearly express your intent; they are not necessarily faster, but they can make your code easier to understand.

```
sumR <- function(x) {  
  total <- 0  
  for (i in seq_along(x)) {  
    total <- total + x[i]  
  }  
  total  
}  
  
cppFunction('double sumC(NumericVector x) {  
  int n = x.size();  
  double total = 0;  
  for(int i = 0; i < n; ++i) {  
    total += x[i];  
  }  
  return total;  
' )
```

Performance

```
> x <- runif(1e6)
> microbenchmark(
+   sum(x),
+   sumC(x),
+   sumR(x)
+ )
```

Unit: microseconds

expr	min	lq	mean	median	uq
sum(x)	801.508	811.3270	892.7033	857.7995	908.4275
sumC(x)	798.635	810.9305	870.4631	849.3895	892.4640
sumR(x)	287038.151	295643.4335	302012.5137	298393.2220	301656.9130
max	neval	cld			
1897.680	100	a			
1256.913	100	a			
369253.282	100	b			

Vector input, vector output

Create a function that computes the Euclidean distance between a value and a vector of values

```
pdistR <- function(x, ys) {  
  sqrt((x - ys) ^ 2)  
}
```

Not obvious that whether `x` is a scalar or not from the function definition. We need to make that clear in the documentation. That is not a problem in the `C++` version because we have to be explicit about types:

```
cppFunction('NumericVector pdistC(double x, NumericVector ys) {  
  int n = ys.size();  
  NumericVector out(n);  
  
  for(int i = 0; i < n; ++i) {  
    out[i] = sqrt(pow(ys[i] - x, 2.0));  
  }  
  return out;  
'})
```

Performance

```
> y = runif(1e6)
> microbenchmark(pdistr(0,y),pdisc(0,y),times=10000L)
Unit: milliseconds
      expr      min       lq      mean    median       uq      max
pdistr(0, y) 5.442920 6.218503 7.393230 6.511349 7.735468 99.9601
pdisc(0, y) 4.101499 4.982845 6.002427 5.177412 6.155853 95.7577
neval cld
10000  b
10000  a
```

Note that because the **R** version is fully vectorised, it is already going to be fast. Assuming it took you 10 minutes to write the **C++** function, you need to run it **600,000** times to make rewriting worthwhile.

The reason why the **C++** function is faster is subtle, and relates to memory management. The **R** version needs to create an intermediate vector the same length as **y** (**x - ys**), and allocating memory is an expensive operation. The **C++** function avoids this overhead because it uses an intermediate scalar.

Matrix input, vector output

Each vector type has a matrix equivalent: `NumericMatrix`, `IntegerMatrix`, `CharacterMatrix`, and `LogicalMatrix`. Using them is straightforward.

```
> cppFunction('NumericVector rowSumsC(NumericMatrix x) {  
+   int nrow = x.nrow(), ncol = x.ncol();  
+       NumericVector out(nrow);  
+       for (int i = 0; i < nrow; i++) {  
+           double total = 0;  
+           for (int j = 0; j < ncol; j++) {  
+               total += x(i, j);  
+           }  
+           out[i] = total;  
+       }  
+       return out;  
+   }')
```

```
> set.seed(2015)
```

```
> x = matrix(sample(100), 10)
```

```
> microbenchmark(rowSums(x), rowSumsC(x), times=10000L)
```

Unit: microseconds

	expr	min	lq	mean	median	uq	max	neval	cld
	rowSums(x)	3.189	3.810	5.025323	4.109	4.458	6060.706	10000	b
	rowSumsC(x)	1.543	1.833	2.410042	1.978	2.156	661.357	10000	a

The main differences: (1) In `C++`, you subset a matrix with `()`, not `[]`. (2) Use `.nrow()` and `.ncol()` methods to get the dimensions of a matrix.

Using sourceCpp()

Inline C++ with `cppFunction()`. This makes presentation simpler, but for real problems, it is usually easier to use stand-alone C++ files and then source them into R using `sourceCpp()`.

This lets you take advantage of text editor support for C++ files (e.g., syntax highlighting) as well as making it easier to identify the line numbers in compilation errors.

Your stand-alone C++ file should have extension `.cpp`, and needs to start with:

```
#include <Rcpp.h>
using namespace Rcpp;
```

Using sourceCpp()

To compile the C++ code, use `sourceCpp("path/to/file.cpp")`. This will create the matching R functions and add them to your current session.

Note that these functions can not be saved in a `.Rdata` file and reloaded in a later session; they must be recreated each time you restart R.

```
#include <Rcpp.h>
using namespace Rcpp;
double meanC(NumericVector x) {
  int n = x.size();
  double total = 0;
  for(int i = 0; i < n; ++i) {
    total += x[i];
  }
  return total / n;
}
```

```
> sourceCpp("~/Dropbox/Umich/Biostat_615_Material/Fall_2015/Rcode/example_mean.cpp")
> x = runif(1e6)
> microbenchmark(meanC(x), mean(x))
```

Unit: microseconds

expr	min	lq	mean	median	uq	max	neval	cld
meanC(x)	798.977	801.792	866.0348	804.8845	869.633	1178.39	100	a
mean(x)	1601.225	1605.066	1755.3828	1638.4170	1710.309	3455.01	100	b

Attributes and other classes

All R objects have attributes, which can be queried and modified with `R.attr()`. Rcpp also provides `.names()` as an alias for the name attribute.

```
> cppFunction('NumericVector attribs() {  
+   NumericVector out = NumericVector::create(1, 2, 3);  
+  
+       out.names() = CharacterVector::create("a", "b", "c");  
+       out.attr("my-attr") = "my-value";  
+       out.attr("class") = "my-class";  
+  
+       return out;  
+   }')
```

```
> attribs()  
a b c  
1 2 3  
attr(,"my-attr")  
[1] "my-value"  
attr(,"class")  
[1] "my-class"
```

Lists and data frames

- Rcpp also provides classes `List` and `DataFrame`, but they are more useful for output than input.
- This is because lists and data frames can contain arbitrary classes but C++ needs to know their classes in advance.
- If the list has known structure, you can extract the components and manually convert them to their C++ equivalents with `as()`.
- For example, the object created by `lm()`, the function that fits a linear model, is a list whose components are always of the same type.

Lists and data frames: Example

- The following code illustrates how you might extract the mean percentage error (`mpe()`) of a linear model.

```
cppFunction('double mpe(List mod) {  
  if (!mod.inherits("lm")) stop("Input must be a linear model");  
  
  NumericVector resid = as<NumericVector>(mod["residuals"]);  
  NumericVector fitted = as<NumericVector>(mod["fitted.values"]);  
  
  int n = resid.size();  
  double err = 0;  
  for(int i = 0; i < n; ++i) {  
    err += resid[i] / (fitted[i] + resid[i]);  
  }  
  return err / n;  
}')
```

```
x = rnorm(1000)  
y = 1+2*x+rnorm(1000)  
fit = lm(y~x)  
mpe(fit)
```

- Note the use of `.inherits()` and the `stop()` to check that the object really is a linear model.

Functions

You can put R functions in an object of type Function. This makes calling an R function from C++ straightforward.

```
> cppFunction(  
+ 'RObject callfun(Function f, int n) {  
+           return f(n);  
+           }'  
+ )  
> microbenchmark(callfun(seq_along,100),seq_along(100))  
Unit: nanoseconds
```

	expr	min	1q	mean
callfun(seq_along, 100)		13964	15229.5	18679.84
seq_along(100)		72	93.0	122.53
median	uq	max	neval	cld
16029	17014.5	79281	100	b
107	149.0	387	100	a

What type of object does an R function return? Use the catchall type `RObject`.