


Biostatistics 615 - Statistical Computing

Lecture 16 Optimizations

Jian Kang

Nov 19, 2015

The Minimization Problem



Specific Objectives

Finding global minimum

- The lowest possible value of the function
- Very hard problem to solve generally

Finding local minimum

- Smallest value within finite neighborhood
- Relatively easier problem


Maximization Problem

- Consider a complex function $f(x)$ (e.g. likelihood)
- Find x which $f(x)$ is maximum or minimum value
- Maximization and minimization are equivalent
 - Replace $f(x)$ with $-f(x)$

Detailed Minimization Strategy

- ① Find 3 points such that
 - $a < b < c$
 - $f(b) < f(a)$ and $f(b) < f(c)$
- ② Then search for minimum by
 - Selecting trial point in the interval
 - Keep minimum and flanking points

Minimization after Bracketing



Step 1: Finding a Bracketing Interval

- Consider two points
 - x-values a, b
 - y-values $f(a) > f(b)$

Bracketing in C++


```
#define SCALE 1.618

void bracket( myFunc foo, double& a, double& b, double& c) {
    double fa = foo(a);
    double fb = foo(b);
    double fc = foo(c = b + SCALE*(b-a) );
    while( fb > fc ) {
        a = b; fa = fb;
        b = c; fb = fc;
        c = b + SCALE * (b-a);
        fc = foo(c);
    }
    // after the loop, fb < fa and fb < fc will hold.
}
```



Step 2 : Finding Minimum After Bracketing

- Given 3 points such that
 - $a < b < c$
 - $f(b) < f(a)$ and $f(b) < f(c)$
- How do we select new trial point?

What is the best location for a new point X ?



What we want



We want to minimize the size of next search interval, which will be either from A to X or from B to C

- If $f(X) < f(B)$, the next search interval will be (B, C)
- If $f(X) > f(B)$, the next search interval will be (A, X)

Minimizing worst case possibility

- Formulae

$$\begin{aligned}w &= \frac{b-a}{c-a} \\z &= \frac{x-b}{c-a}\end{aligned}$$

Segments will have length either $1-w$ or $w+z$.


- Optimal case

$$\begin{cases} 1-w = w+z \\ \frac{z}{1-w} = w \end{cases}$$


- Solve It

$$w = \frac{3-\sqrt{5}}{2} = 0.38197$$


The Golden Search



The Golden Ratio




The Golden Ratio



The number 0.38196 is related to the *golden mean* studied by Pythagoras

The Golden Ratio



- Reduces bracketing by $\sim 40\%$ after function evaluation
- Performance is independent of the function that is being minimized
- In many cases, better schemes are available

Golden Step

```
#define GOLD 0.38196
#define ZEPS 1e-10    // precision tolerance
double goldenStep (double a, double b, double c) {
    double mid = ( a + c ) * .5;
    if ( b > mid )
        return GOLD * (a-b);
    else
        return GOLD * (c-b);
}
```

Golden Search

```
double goldenSearch(myFunc foo, double a, double b, double c, double e) {
    int i = 0;
    double fb = foo(b);
    while ( fabs(c-a) > fabs(b*e) ) {
        double x = b + goldenStep(a, b, c);
        double fx = foo(x);
        if ( fx < fb ) {
            (x > b) ? ( a = b ) : ( c = b);
            b = x; fb = fx;
        }
        else {
            (x < b) ? ( a = x ) : ( c = x );
        }
        ++i;
    }
    std::cout << "i = " << i << ", b = " << b << ", f(b) = " << foo(b) << std::endl;
    return b;
}
```

A running example

Finding minimum of $f(x) = -\cos(x)$

```
class myFunc {
public:
    double operator() (double x) const {
        return 0-cos(x);
    }
};
..
int main(int argc, char** argv) {
    myFunc foo;
    goldenSearch(foo,0-M_PI/4,M_PI/4,M_PI/2,1e-5);
    return 0;
}
```

Results

i = 66, b = -4.42163e-09, f(b) = -1


R example of minimization

```
> optimize(cos,interval=c(0-pi/4,pi/2),maximum=TRUE)
$maximum
[1] -8.648147e-07

$objective
[1] 1
```

- Parabola Method: Using a quadratic approximation of the function may achieve better optimization results; Likely more efficient reduction, but not always guaranteed.
- Brent's Method: Combination of above two methods. More efficient than both.
- The Newton-Raphson Method: Quadratic convergence. Requires derivatives.

Multidimensional Optimization : A mixture distribution



A general mixture distribution

$$p(x; \pi, \phi, \eta) = \sum_{i=1}^k \pi_i f(x; \phi_i, \eta)$$

x observed data

π mixture proportion of each component

f the probability density function

ϕ parameters specific to each component

η parameters shared among components

k number of mixture components

Problem : Maximum Likelihood Estimation

Finding Maximum-likelihood

Find parameters that maximizes the likelihood of the entire sample

$$L = \prod_i p(x_i | \pi, \phi, \eta)$$

Calculating in log-space

Or equivalently, consider log-likelihood to avoid underflow

$$l = \sum_i \log p(x_i | \pi, \phi, \eta)$$

Gaussian MLE in single-dimensional space

$$p(x; \mu, \sigma^2) = \mathcal{N}(x; \mu, \sigma^2)$$

Given x , what is the MLE parameters of μ and σ^2 ?

- Analytical solution does exist
- $\hat{\mu} = \sum_{i=1}^n x_i / n$
- $\hat{\sigma}^2 = \sum_{i=1}^n (x_i - \hat{\mu})^2 / n$

MLE in Gaussian mixture

Parameter estimation in Gaussian mixture

- No analytical solution
- Numerical optimization required
- Multi-dimensional optimization problem
 - $\pi_1, \pi_2, \dots, \pi_k$
 - $\mu_1, \mu_2, \dots, \mu_k$
 - $\sigma_1^2, \sigma_2^2, \dots, \sigma_k^2$

Possible methods for multi-dimensional optimization

- Nelder-Mead Method
- Gradient Descent
- Newton's Method
- Expectation Maximization
- Simulated Annealing
- Markov-Chain Monte Carlo

The Nelder-Mead Method

- a.k.a. downhill simplex method or amoeba method
- Calculate likelihoods at simplex vertexes
 - Geometric shape with $k + 1$ corners
 - A triangle in $k = 2$ dimensions
 - A tetrahedron in $k = 3$ dimensions
- Simplex *crawls*
 - Towards minimum
 - Away from maximum
- A commonly used nonlinear optimization method, without using derivatives.
- May converge to non-stationary points.

(http://en.wikipedia.org/wiki/Nelder%E2%80%93Mead_method)


How the Nelder-Mead Method Works

(<http://userpages.umbc.edu/~rostamia/2013-01-math625/images/nelder-mead.gif>)


Nelder-Mead Method in Two Dimensions

- Evaluate functions at three vertexes
 - The highest (worst) point
 - The next highest point
 - The lowest (best) point
- Intuition
 - Move away from high point, towards low point


Direction for Optimization




Reflection




Reflection and Expansion



Contraction (1-dimension)




Multiple Contraction



"passing through the
eye of a needle"

If a simple contraction doesn't
improve things, then try moving
all points towards the current
minimum

Summary : The Nelder-Mead Method (in 3-D)



R function for optimization

General purpose optimizations

```
optim(par, fn, gr = NULL, ...,  
      method = c("Nelder-Mead", "BFGS", "CG", "L-BFGS-B", "SANN",  
                  "Brent"),  
      lower = -Inf, upper = Inf,  
      control = list(), hessian = FALSE)
```

Example,

```
> fn = function(x)  
+ return(100*(x[2]-x[1]^2)^2 + (1-x[1])^2)  
> results = optim(par=c(-1.2,1),fn=fn,method="Nelder-Mead")  
> results$par  
[1] 1.000260 1.000506  
> results$value  
[1] 8.825241e-08
```

Implementing the Nelder-Mead Method in C++

```
template <class F> // F is a function object
class simplex615 { // contains (dim+1) points of size (dim)
protected:
    std::vector<std::vector<double> > X; // (dim+1)*dim matrix
    std::vector<double> Y;                // (dim+1) vector
    std::vector<double> midPoint;          // variables for update
    std::vector<double> thruLine;          // variables for update
    int dim, idxLo, idxHi, idxNextHi;      // dimension, min, max, 2ndmax values
    void evaluateFunction(F& foo); // evaluate function value at each point
    void evaluateExtremes();           // determine the min, max, 2ndmax
    void prepareUpdate();              // calculate midPoint, thruLine
    bool updateSimplex(F& foo, double scale); // for reflection/expansion..
    void contractSimplex(F& foo); // for multiple contraction
    static int check_tol(double fmax, double fmin, double ftol); // check tolerance
public:
    simplex615(double* p, int d);        // constructor with initial points
    void amoeba(F& foo, double tol); // main function for optimization
    std::vector<double>& xmin();           // optimal x value
    double ymin();                        // optimal y value
};
```

- Data representation
 - Each $x[i]$ is point of the simplex
 - $y[i]$ corresponds to $f(X[i])$
 - `midPoint` is the average of all points (except for the worst point)
 - `thruLine` is vector from the worst point to the `midPoint`

- Data representation
 - Each $x[i]$ is point of the simplex
 - $y[i]$ corresponds to $f(X[i])$
 - `midPoint` is the average of all points (except for the worst point)
 - `thruLine` is vector from the worst point to the `midPoint`
- Reflection, Expansion and Contraction
After calculating `midPoint` and `thruLine`
 - Reflection Call `updateSimplex(foo, -1.0)`
 - Expansion Call `updateSimplex(foo, -2.0)`
 - Contraction Call `updateSimplex(foo, 0.5)`

Initializing a Simplex

```
// constructor of simplex615 class : initial point is given
template <class F>
simplex615<F>::simplex615(double* p, int d) : dim(d) { // set dimension
    // Determine the space required
    X.resize(dim+1);        // X is vector-of-vector, like 2-D array
    Y.resize(dim+1);        // Y is function value at each simplex point
    midPoint.resize(dim);
    thruLine.resize(dim);
    for(int i=0; i < dim+1; ++i) {
        X[i].resize(dim); // allocate the size of content in the 2-D array
    }
    // Initially, make every point in the simplex identical
    for(int i=0; i < dim+1; ++i)
        for(int j=0; j < dim; ++j)
            X[i][j] = p[j]; // set each simple point to the starting point
    // then increase each dimension by one unit except for the last point
    for(int i=0; i < dim; ++i)
        X[i][i] += 1.;      // this will generate a simplex
}
```


Evaluating function values at each simplex point

```
// simple function for evaluating the function value at each simple point
// after calling this function Y[i] = foo(X[i]) should hold
template <class F>
void simplex615<F>::evaluateFunction(F& foo) {
    for(int i=0; i < dim+1; ++i) {
        Y[i] = foo(X[i]); // foo is a function object, which will be visited later
    }
}
```

Determine the best, worst, and the second-worst points

```
template <class F>
void simplex615<F>::evaluateExtremes() {
    if ( Y[0] > Y[1] ) { // compare the first two points
        idxHi = 0; idxLo = idxNextHi = 1;
    }
    else {
        idxHi = 1; idxLo = idxNextHi = 0;
    }
    // for each of the next points
    for(int i=2; i < dim+1; ++i) {
        if ( Y[i] <= Y[idxLo] ) // update the best point if lower
            idxLo = i;
        else if ( Y[i] > Y[idxHi] ) { // update the worst point if higher
            idxNextHi = idxHi; idxHi = i;
        }
        else if ( Y[i] > Y[idxNextHi] ) // update also if it is the 2nd-worst point
            idxNextHi = i;
    }
}
```

Direction for Optimization




Determining the direction for optimization

```
template <class F>
void simplex615<F>::prepareUpdate() {
    for(int j=0; j < dim; ++j) {
        midPoint[j] = 0;           // average of all points but the worst point
    }
    for(int i=0; i < dim+1; ++i) {
        if ( i != idxHi ) {        // exclude the worst point
            for(int j=0; j < dim; ++j) {
                midPoint[j] += X[i][j];
            }
        }
    }
    for(int j=0; j < dim; ++j) {
        midPoint[j] /= dim;        // take average
        thruLine[j] = X[idxHi][j] - midPoint[j]; // direction for optimization
    }
}
```


Updating simplex along the line

```
// scale determines which point to evaluate along the line
// scale = 1 : worse point, scale = 0 : midPoint
template <class F>
bool simplex615<F>::updateSimplex(F& foo, double scale) {
    std::vector<double> nextPoint;    // next point to evaluate
    nextPoint.resize(dim);
    for(int i=0; i < dim; ++i) {
        nextPoint[i] = midPoint[i] + scale * thruLine[i];
    }
    double fNext = foo(nextPoint);
    if ( fNext < Y[idxHi] ) {          // update only maximum values (if possible)
        for(int i=0; i < dim; ++i) { // because the order can be changed with
            X[idxHi][i] = nextPoint[i]; // evaluateExtremes() later
        }
        Y[idxHi] = fNext;
        return true;
    }
    else {
        return false;                // never mind if worse than the worst
    }
}
```


Reflection




Reflection and Expansion



Contraction (1-dimension)



Multiple Contraction



"passing through the
eye of a needle"

If a simple contraction doesn't
improve things, then try moving
all points towards the current
minimum

Updating simplex along the line

```
// if none of the tried points make things better
// reduce the search space towards the minimum point
template <class F>
void simplex615<F>::contractSimplex(F& foo) {
    for(int i=0; i < dim+1; ++i) {
        if ( i != idxLo ) { // except for the minimum point
            for(int j=0; j < dim; ++j) {
                X[i][j] = 0.5*( X[idxLo][j] + X[i][j] ); // move the point towards minimum
            }
            Y[i] = foo(X[i]); // re-evaluate the function
        }
    }
}
```

Putting things together

```
template <class F>
void simplex615<F>::amoeba(F& foo, double tol) {
    evaluateFunction(foo); // evaluate the function at the initial points
    while(true) {
        evaluateExtremes(); // determine three important points
        prepareUpdate(); // determine direction for optimization

        if ( check_tol(Y[idxHi],Y[idxLo],tol) ) break; // check convergence
        updateSimplex(foo, -1.0); // reflection
        if ( Y[idxHi] < Y[idxLo] ) {
            updateSimplex(foo, -2.0); // expansion
        }
        else if ( Y[idxHi] >= Y[idxNextHi] ) {
            if ( !updateSimplex(foo, 0.5) ) { // 1-d contraction
                contractSimplex(foo); // multiple contractions
            }
        }
    }
}
```

- A general purpose minimization routine
 - Works in multiple dimensions
 - Uses only function evaluations
 - Does not require derivatives

Checking convergence

```
// Note that the function is declared as "static" function as
//
// static int check_tol(double fmax, double fmin, double ftol);
//
// because it does not use any member variables
template <class F>
int simplex615<F>::check_tol(double fmax, double fmin, double ftol) {
    // calculate the difference
    double delta = fabs(fmax - fmin);
    // calculate the relative tolerance
    double accuracy = (fabs(fmax) + fabs(fmin)) * ftol;
    // check if difference is within tolerance
    return (delta < (accuracy + ZEPS));
}
```

Using the Nelder-Mead Method Implementation

```
#include <vector>
#include <cmath>
#include <iostream>
#include "simplex615.h"
#define ZEPS 1e-10

int main(int main, char** argv) {
    double point[2] = {-1.2, 1}; // initial point to start

    arbitraryFunc foo;           // WILL BE DISCUSSED LATER
    simplex615<arbitraryFunc> simplex(point, 2); // create a simplex
    simplex.amoeba(foo, 1e-7);    // optimize for a function
    // print outputs
    std::cout << "Minimum = " << simplex.ymin() << ", at ("
                << simplex.xmin()[0] << ", " << simplex.xmin()[1]
                << ")" << std::endl;
    return 0;
}
```

Defining arbitraryFunc

```
// function object used as an argument
class arbitraryFunc {
public:
    double operator() (std::vector<double>& x) {
        //  $f(x_0, x_1) = 100 \cdot (x_1 - x_0^2)^2 + (1 - x_0)^2$ 
        return 100*(x[1]-x[0]*x[0])*(x[1]-x[0]*x[0])+(1-x[0])*(1-x[0]);
    }
};
```


A working example

```
Minimum = 1.35567e-11, at (0.999999, 0.999997)
```

Normal Density

Normal density function

$$f(x|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp \left[-\frac{1}{2} \left(\frac{x-\mu}{\sigma} \right)^2 \right]$$

Implementation

```
class NormMix615 {  
public:  
    static double dnorm(double x, double mu, double sigma) {  
        return 1.0 / (sigma * sqrt(M_PI * 2.0)) *  
            exp (-0.5 * (x - mu) * (x-mu) / sigma / sigma);  
    }  
    ...  
};
```

Gaussian mixture distribution

Density function

$$p(x|k, \pi, \mu, \sigma) = \sum_{i=1}^k \pi_i f(x|\mu_i, \sigma_i)$$

Implementation (within NormMix615)

```
static double dmix(double x, std::vector<double>& pis,
                  std::vector<double>& means, std::vector<double>& sigmas) {
    double density = 0;
    for(int i=0; i < (int)pis.size(); ++i)
        density += pis[i] * dnorm(x, means[i], sigmas[i]);
    return density;
}
```

Likelihood of multiple observations

Calculating in log-space

$$L = \prod_i p(x_i | \pi, \mu, \sigma)$$
$$l = \sum_i \log p(x_i | \pi, \mu, \sigma)$$

Implementation (within NormMix615)

```
static double mixLLK(std::vector<double>& xs, std::vector<double>& pis,
                    std::vector<double>& means, std::vector<double>& sigmas) {
    int i=0;
    double llk = 0.0;
    for(int i=0; i < xs.size(); ++i)
        llk += log(dmix(xs[i], pis, means, sigmas));
    return llk;
}
```

Gaussian Mixture Function Object

```
class NormMix615 {
public:  // these are internal function
    static double dnorm(double x, double mu, double sigma);
    static double dmix(...);
    static double mixLLK(...);
};

class LLKNormMixFunc {
public:  // below are public functions
    LLKNormMixFunc(int k, std::vector<double>& y) :
        numComponents(k), data(y), numFunctionCalls(0) {}
    // core function - called when foo() is used
    // x is the combined list of MLE parameters (pis, means, sigmas)
    double operator() (std::vector<double>& x);
    void assignPriors(std::vector<double>& x, std::vector<double>& priors);
    std::vector<double> data;
    int numComponents;
    int numFunctionCalls;
};
```

Avoiding boundary conditions

Problem

- The simplex algorithm do not know that $0 \leq \pi_i \leq 1$, and $\sum_{i=1}^n \pi_i = 1$
- During the iteration of simplex algorithm, it is possible that π_i goes out of bound

Possible solutions

- Modify simplex algorithm to avoid boundary conditions
- Transform the parameter space to infinite ranges

Transforming the parameter space

Constraints

- $0 \leq \pi_i \leq 1$
- $\sum_{i=1}^n \pi_i = 1$

Mapping between the space

- Given $x \in \mathbb{R}^{n-1}$, for $i = 1, \dots, n-1$
- $\pi_i = \frac{1}{1+e^{-x_i}} (1 - \sum_{j=1}^{i-1} \pi_j)$
- $\pi_n = 1 - \sum_{i=1}^{n-1} \pi_i.$

Q: how about σ ?

Implementing likelihood of data

```
double LLKNormMixFunc::operator() (std::vector<double>& x) { // x has (3*k-1) dims
    std::vector<double> priors;
    std::vector<double> means;
    std::vector<double> sigmas;
    // transform (k-1) real numbers to priors
    assignPriors(x, priors);
    for(int i=0; i < numComponents; ++i) {
        means.push_back(x[numComponents-1+i]);
        sigmas.push_back(x[2*numComponents-1+i]);
    }
    return 0-NormMix615::mixLLK(data, priors, means, sigmas);
}
```



Transforming between bounded and unbounded space of priors

```
void LLKNormMixFunc::assignPriors(std::vector<double>& x, std::vector<double>& priors) {
    priors.clear();
    // convert priors (from [k-1]-d real scale to [k]-d simplex scale)
    double p = 1.;
    for(int i=0; i < numComponents-1; ++i) {
        double logit = 1./(1.+exp(0-x[i]));
        priors.push_back(p*logit);
        p = p*(1.-logit);
    }
    priors.push_back(p);
}
```

Nelder-Mead Method for Gaussian Mixture

```
#include <iostream>
#include <fstream>
#include "simplex615.h"
#include "normMix615.h"
#include "llkNormMixFunc.h"
#define ZEPS 1e-10
int main(int main, char** argv) {
    double point[5] = {0, -1, 1, 1, 1}; // 50:50 mixture of N(-1,1) and N(1,1)
    simplex615<LLKNormMixFunc> simplex(point, 5);
    std::vector<double> data;           // input data
    std::ifstream file(argv[1]);        // open file
    double tok;                         // temporary variable
    while(file >> tok) data.push_back(tok); // read data from file
    LLKNormMixFunc foo(2, data);        // 2-dimensional mixture model
    simplex.amoeba(foo, 1e-7);           // run the Nelder-Mead Method
    std::cout << "Minimum = " << simplex.ymin() << ", at pi = "
              << (1./(1.+exp(0-simples.xmin()[0]))) << ", " << "between N("
              << simplex.xmin()[1] << ", " << simplex.xmin()[3] << ") and N("
              << simplex.xmin()[2] << ", " << simplex.xmin()[4] << ")" << std::endl;
    return 0;
}
```

A working example



A working example

Simulation of data

```
> x <- rnorm(1000)
> y <- rnorm(500)+5
> write.table(matrix(c(x,y),1500,1), 'mix.dat', row.names=F, col.names=F)
```

A Running Example

Minimum = 3043.46, at $\pi = 0.667271$,
between $N(-0.0304604, 1.00326)$ and $N(5.01226, 0.956009)$