

Chapter 11: Shrinkage Methods

Stats 500, Fall 2015
Brian Thelen, University of Michigan
443 West Hall, bjthelen@umich.edu

Shrinkage Methods


- Principal components regression
- Partial least squares
- Ridge regression

Principal Components (PC)

Main uses:

- **Reduce the dimensionality of the data**
- Find linear combinations of predictors that explain the most variation in the data
- Facilitate visualization
- In regression, makes predictors orthogonal to each other

PC Illustration



Definition of PCs

- 1 Center each variable by its mean,
$$x_j - \bar{x}_j \Rightarrow X_{n \times p}.$$
- 2 Find u_1 such that $\text{var}(Xu_1)$ is maximized subject to $u_1^T u_1 = 1$.
- 3 For each $k = 2, \dots, p$, given u_1, \dots, u_{k-1} , find u_k such that $\text{var}(Xu_k)$ is maximized subject to $u_k^T u_k = 1$ and $u_k^T u_j = 0$ for all $j = 1, \dots, k-1$.
- 4 Vectors u_j are called the PC directions.
- 5 Vectors $z_j = Xu_j$ are called the principal components of X .

Remarks

- $z_j = Xu_j$ are **projections** of data points on the direction u_j .
- u_j are the **eigenvectors** of $X^T X$.
- $\text{var}(Xu_j) = \lambda_j$, the **eigenvalues** of $X^T X$.
- $\lambda_1 \geq \lambda_2 \geq \dots \lambda_p$
- Combine vectors into a matrix and get

$$\text{diag}(\lambda_1, \dots, \lambda_p) = Z^T Z = U^T X^T X U$$

- Recommended: standardize each variable first.

Principal Components Regression

PCR replaces the regression model

$$y = \beta_0 + \beta_1 x_1 + \dots \beta_p x_p$$

with

$$y = \beta'_0 + \beta'_1 \mathbf{z}_1 + \dots + \beta'_k \mathbf{z}_k$$

Note: PCs are centered, so $\hat{\beta}'_0 = \bar{y}$.

How do we pick the number of PCs?

- Typically most variation in X can be represented by a few principal components – **Dimension reduction**.
- Can choose k to explain certain **percent of variation**: pick first k so that

$$\sum_{i=1}^k \lambda_i \geq (1 - \alpha) \sum_{i=1}^p \lambda_i$$

- Can look at the **scree plot** (sq. rooted eigenvalues in decreasing order) and look for a gap in eigenvalues
- More sophisticated methods for estimating **intrinsic dimension** of the data

Food Analyzer Example

- Response: fat content
- Predictors: 100 channel spectrum of absorbances
- Number of data points: $n = 215$
- Number of predictors: $p = 100$

Prediction Performance

Goal : build a model that predicts well on **future data** .

- Divide the data into two groups: **training data** and **test data** .
- Build the models using the training data and evaluate them on the test data.

Food Analyzer Example Continued

```
> library(faraway)
> data(meatspec)
> dim(meatspec)
[1] 215 101
## Training and test data
> trainmeat = meatspec[1:172,]
> testmeat = meatspec[173:215,]

## Linear model
> modlm = lm(fat ~ ., tr)
> summary(modlm)$r.squared
[1] 0.9970196

## Root mean squared error
> rmse = function(x, y) { sqrt(mean( (x - y)^2 ))}
> rmse(modlm$fit, trainmeat$fat)
[1] 0.6903167
> ## Prediction
> rmse( predict(modlm, newdata=testmeat), testmeat$fat )
[1] 3.814000
```

```
## AIC
> modlm2 = step(modlm)
> rmse( modlm$fit, tr$fat )
[1] 0.7095069
> rmse( predict(modlm2, newdata=testmeat), testmeat$fat )
[1] 3.590245


## Principal components regression
> library(stats)
> meatpca = prcomp(tr[, -101])
## Square root of the eigenvalues
> round(meatpca$sdev, 3)
[1] 5.055 0.511 0.282 0.168 0.038 0.025 0.014
[8] 0.011 0.005 0.003 0.002 0.002 0.001 0.001
... ..
> matplot(1:100, meatpca$rot[, 1:3], type="l",
          xlab="Frequency", ylab="", lwd=3)
```

```


## Make a scree plot (to choose number of PCs k)
> plot(1:10, meatpca$sdev[1:10], type="l",
      xlab="PC number", ylab="SD of PC")
## Fit all PCRs at once and calculate test RMSE for each k
> library(pls)
> modpcr = pcr(fat ~ ., data=tr, ncomp=50)
> rmsmeat = NULL
> for (k in 1:50) {
+   pv = predict(modprc, newdata=testmeat, ncomp=k)
+   rmsmeat[k] = rmse(pv, te$fat ) }
> plot(rmsmeat, xlab="PC number", ylab="Test RMS")
# scree plot suggestion
> rmsmeat[5]
[1] 3.533628
> which.min(rmsmeat)
[1] 27
> rmsmeat[27]
[1] 1.854858

```

Comparison of Coefficients




Coefficients from Linear Model




Effective coefficients from
Pr. Component Regression (n=5)

```
plot(modlm$coef[-1],xlab="Frequency",ylab="Coefficient",  
+ type="l")  
coefplot(modpcr, ncomp=5, xlab="Frequency",main="")
```


Food Analyzer: First 3 PCs



Food Analyzer: Scree plot



Food Analyzer: Test RMSE



Cross-Validation

- We cannot use the test data to pick k .
- Setting some of the training data aside for validation is possible, but reduces training sample size

A possible solution: **K -fold cross-validation** :


- ① Divide the data into K parts
- ② Use parts $2, \dots, K$ as *training* data and part 1 as *test* data. Compute prediction errors on part 1.
- ③ Repeat for each part
- ④ Average prediction errors from all parts and pick k minimizing the average.

Food Analyzer Example: Cross-validation

```
> modpcr2 = pcr(fat ~ ., data=trainmeat, ncomp=50,  
> +           validation="CV", segments = 10)  
> rmsCV= RMSEP(modpcr2, estimate='CV')  
> which.min(rmsCV$val)  
19  
  
# Another try  
> modpcr2 = pcr(fat ~ ., data=trainmeat, ncomp=50,  
> +           validation="CV", segments = 10)  
> rmsCV= RMSEP(modpcr2, estimate='CV')  
> which.min(rmsCV$val)  
21  
  
## Plot the RMSE; k=0 is the model with intercept only  
> plot(rmsCV$val, xlab="PC number", ylab="CV RMS")  
## Get test error  
> yfit = predict(modpcr2, newdata=testmeat, ncomp=21)  
> rmse(testmeat$fat, yfit)  
[1] 2.214545
```

CV tends to underestimate the real test RMSE – often close

Food Analyzer Example (PCR)



Partial Least Squares

- **PCR** ignores y when building z 's
- Partial least squares (**PLS**) chooses z 's that are best at predicting y .
- PLS can handle multiple-output regression (vector-valued y)
- PLS does not solve a well-defined modelling problem
- Many algorithms for PLS exist
- Also need to select number of components
- No interpretation

Partial Least Squares Continued

Algorithm:

- 1 Center y , center and standardize each x_j
- 2 Regress y on each x_j *separately* to get α_j
- 3 Construct $z_1 = \sum \alpha_j x_j$, which is the first PLS component
- 4 Regress y on z_1 to get $\hat{\beta}_1$
- 5 Orthogonalize each x_j with respect to z_1
- 6 Continue until the final model is fit:


$$\hat{y} = \bar{y} + \hat{\beta}_1 z_1 + \cdots \hat{\beta}_k z_k$$

Food Analyzer Example

```
> ## Partial least squares
> modpls = plsr(fat ~ ., data=tr, ncomp=50, validation="CV")
> # plot RMSE estimated by CV
> pls_rmsCV = RMSEP(modpls, estimate='CV')
> plot(pls_rmsCV$val, xlab="PC number", ylab="CV RMS")
> which.min(pls_rmsCV$val)
[1] 14
> ## RMSE on the training data
> dim(modpls$fit)
[1] 172    1   50
> rmse(modpls$fit[, , 14], tr$fat)
[1] 1.952796

> ## RMSE on the test data
> ypred.test = predict(plsg, newdata=testmeat)
> dim(ypred.test)
[1] 43    1   50
> rmse(ypred.test[, , 14], testmeat$fat)
[1] 2.011180
```

Food Analyzer Example (PLS)



Ridge Regression

When to use this: the predictors are collinear and the usual least squares estimates are unstable.

Center y , center and standardize each x_j . Consider

$$\min (y - X\beta)^T (y - X\beta) + \lambda \|\beta\|^2$$

β does not include the intercept β_0 . Solution:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

Remarks


- $\lambda = 0$ reduces to the usual least squares solution
- $\lambda \rightarrow \infty$ implies $\hat{\beta} \rightarrow 0$
- Alternative formulation: solve

$$\min (y - X\beta)^T (y - X\beta)$$

subject to $\|\beta\|^2 \leq C$.

- Another shrinkage method is **Lasso** : replace $\|\beta\|^2 = \sum_i \beta_i^2$ with $\sum_i |\beta_i|$.
- Lasso shrinks many coefficients to exactly 0; ridge makes all coefficients smaller but does not set them to 0.

Picture: Ridge Regression and Lasso



Bias-Variance Trade-off

Why does ridge regression work? Is it biased?

$$E(\hat{\beta}) = (X^T X + \lambda I)^{-1}(X^T X)\beta$$

Goal: mean squared error

$$\begin{aligned} \mathbf{E}(\hat{\mathbf{y}} - \mathbf{E}(\mathbf{y}))^2 &= (E(\hat{y}) - E(y))^2 + E(\hat{y} - E(\hat{y}))^2 \\ &= \text{Bias}^2 + \text{Variance} \end{aligned}$$

Sometimes introducing a small bias leads to a large drop in variance.

Food Analyzer: Ridge Regression

```
> library(MASS)
> ## The function will center the training data
> modridge = lm.ridge(fat ~ ., lambda=seq(0, 5e-8, 1e-9),
+ data = trainmeat)
> matplot(modridge$lambda, t(modridge$coef), type="l", lty=1,
+ xlab=expression(lambda), ylab=expression(hat(beta)) )


> ## Select an appropriate lambda
> select(modridge)
modified HKB estimator is 1.058342e-08
modified L-W estimator is 0.7096864
smallest value of GCV at 1.8e-08
> abline(v=1.8e-8)

> ## No fitted values returned - compute yourself
> yfit = modridge$ym + scale(tr[,-101], center=gridge$xm,
+ scale=modridge$scales ) %*% modridge$coef[, 19]
> # RMSE on training data
> rmse(yfit, trainmeat$fat)
[1] 0.8045392
```

```
## No predict() function for ridge
## Predict for test data, but need to center with training means
> ypred = modridge$ym + scale(testmeat[, -101], center=
+ modridge$xm, scale = modridge$scales ) %*% modridge$coef[, 19]
> rmse(ypred, testmeat$fat)
[1] 4.096579
```

```
## One really bad prediction
> c( ypred[13], testmeat$fat[13] )
      185
11.18769 34.80000
> rmse( ypred[-13], testmeat$fat[-13] )
[1] 1.976548
```

Food Analyzer Example Continued



Summary

- Main reason to use shrinkage: too many predictors or collinearity
- Interpretation is usually lost
- Ridge is still a linear model in the original predictors but no inference
- Prediction is usually improved by shrinkage
- All require selecting a **tuning parameter** (number of components for PCR and PLS, λ for ridge) – need validation data or cross-validation.